

Introduction to Seaborn

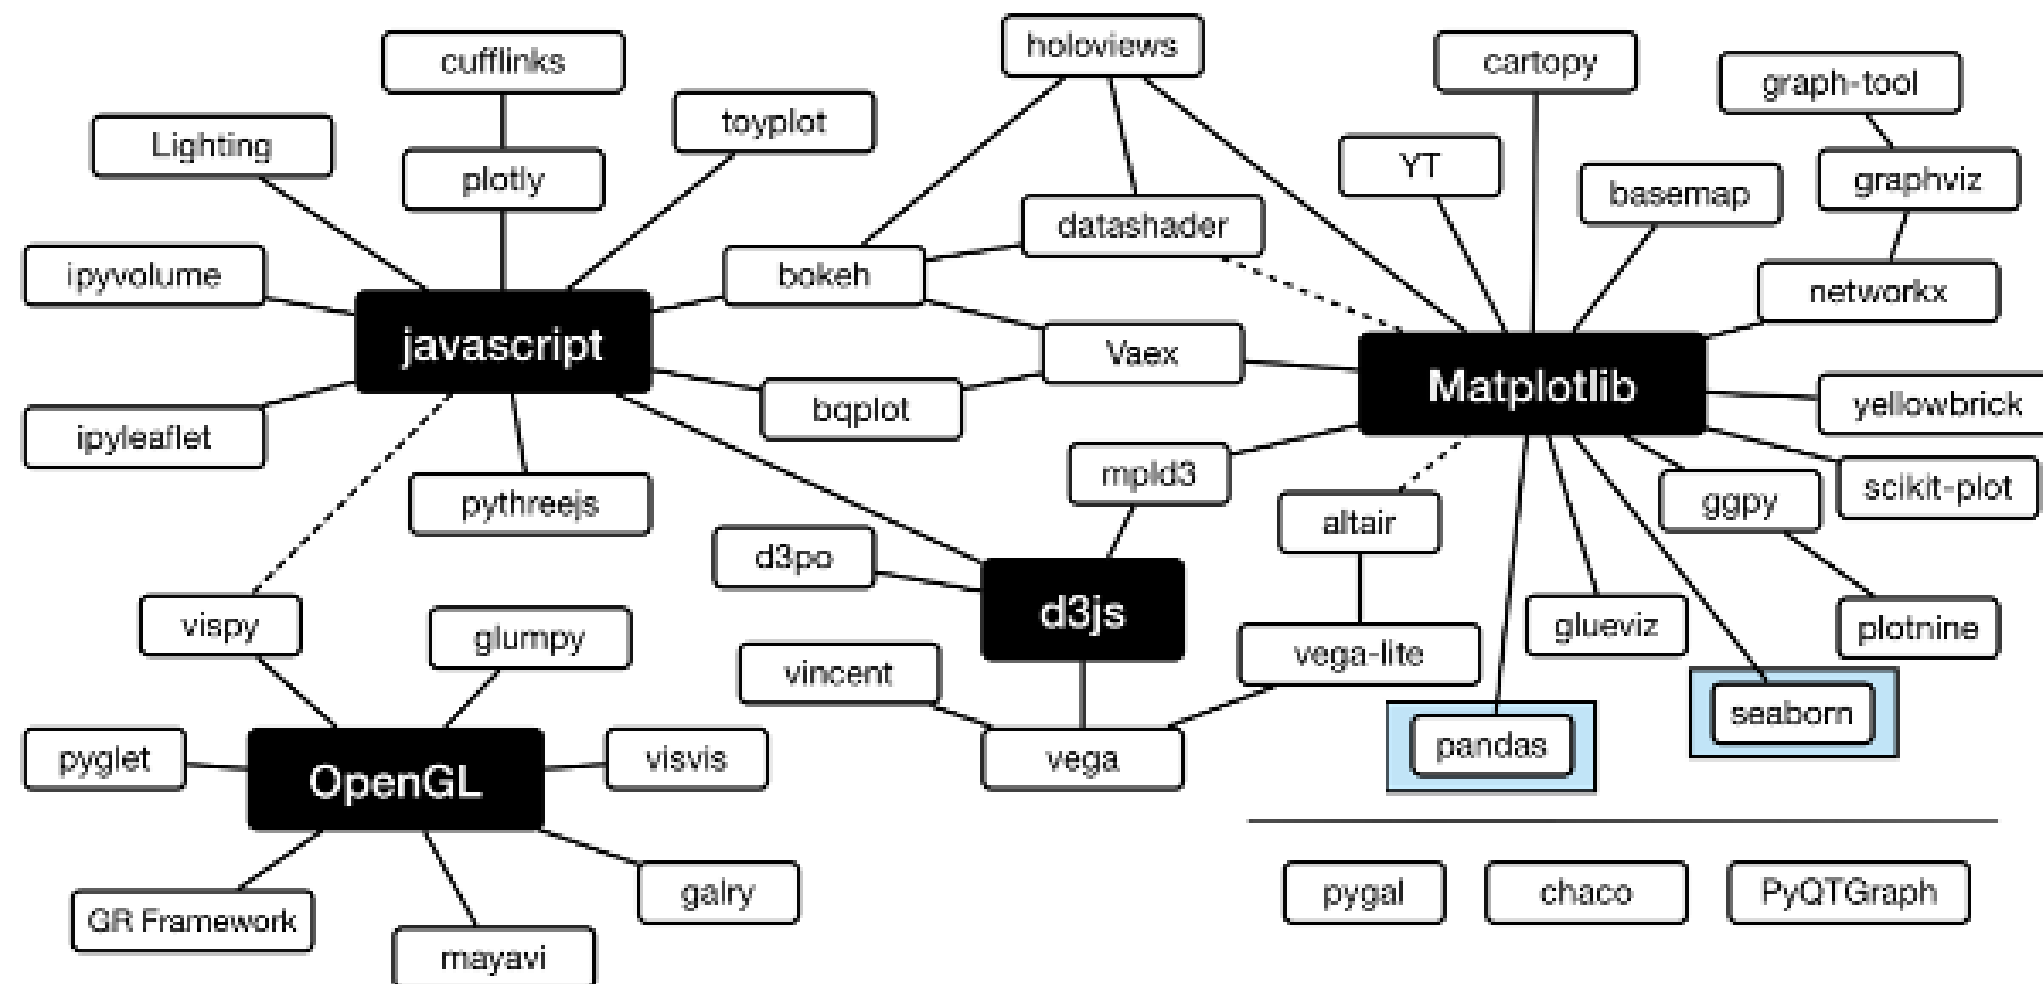
INTERMEDIATE DATA VISUALIZATION WITH SEABORN



Chris Moffitt
Instructor

Python Visualization Landscape

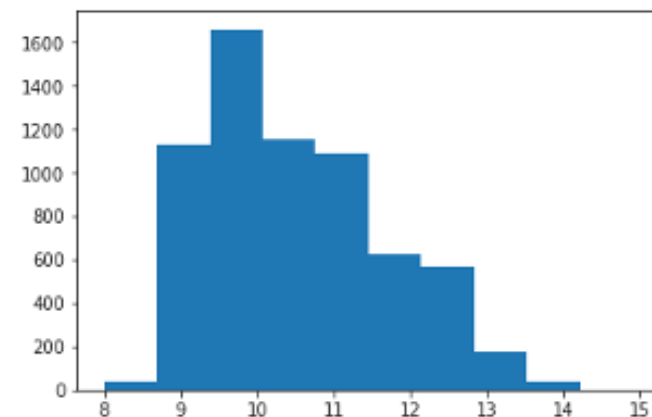
- The python visualization landscape is complex and can be overwhelming



Matplotlib

- `matplotlib` provides the raw building blocks for Seaborn's visualizations
- It can also be used on its own to plot data


```
import matplotlib.pyplot as plt
import pandas as pd
df = pd.read_csv("wines.csv")
fig, ax = plt.subplots()
ax.hist(df['alcohol'])
```



Pandas

- `pandas` is a foundational library for analyzing data
- It also supports basic plotting capability

```
import pandas as pd
df = pd.read_csv("wines.csv")
df['alcohol'].plot.hist()
```




Seaborn

- Seaborn supports complex visualizations of data
- It is built on matplotlib and works best with pandas' dataframes

Seaborn

- The `distplot` is similar to the histogram shown in previous examples
- By default, generates a Gaussian Kernel Density Estimate (KDE)

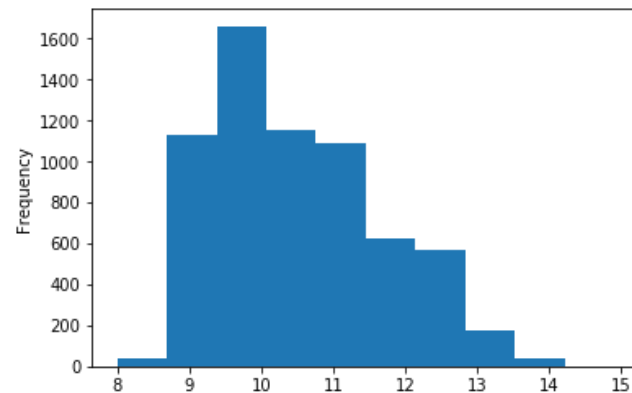
```
import seaborn as sns
sns.distplot(df['alcohol'])
```



Histogram vs. Distplot

- Pandas histogram

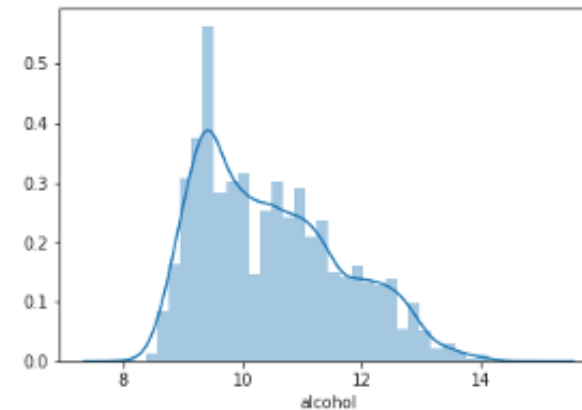
```
df['alcohol'].plot.hist()
```



- Actual frequency of observations
- No automatic labels
- Wide bins

- Seaborn distplot

```
sns.distplot(df['alcohol'])
```



- Automatic label on x axis
- Muted color palette
- KDE plot
- Narrow bins

Let's practice!

INTERMEDIATE DATA VISUALIZATION WITH SEABORN

Using the distribution plot

INTERMEDIATE DATA VISUALIZATION WITH SEABORN

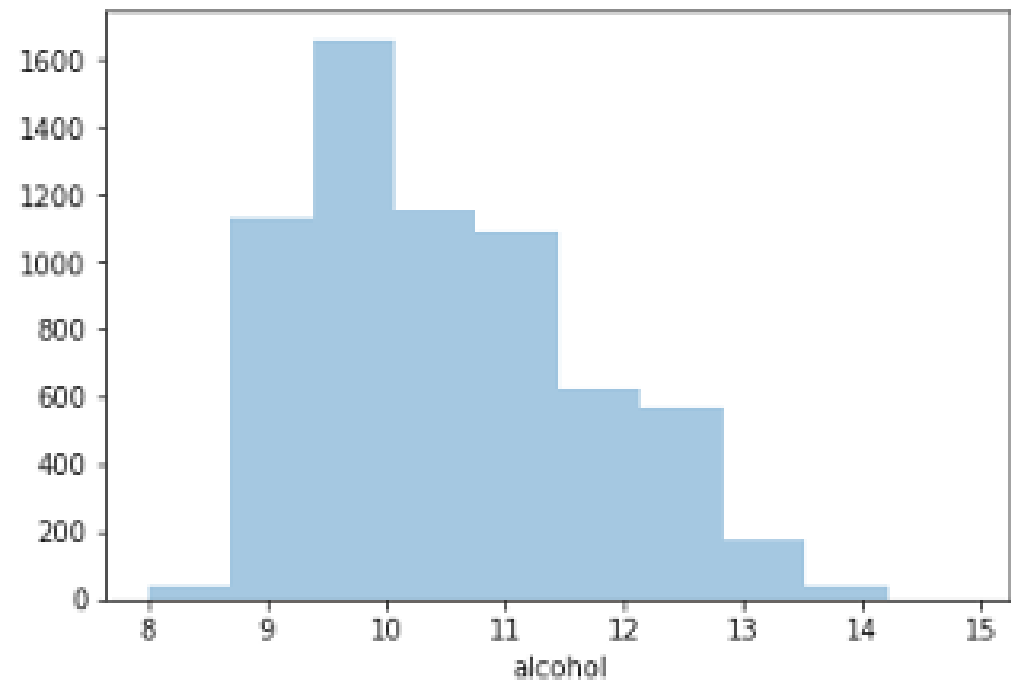


Chris Moffitt
Instructor

Creating a histogram

- Distplot function has multiple optional arguments
- In order to plot a simple histogram, you can disable the kde and specify the number of bins to use

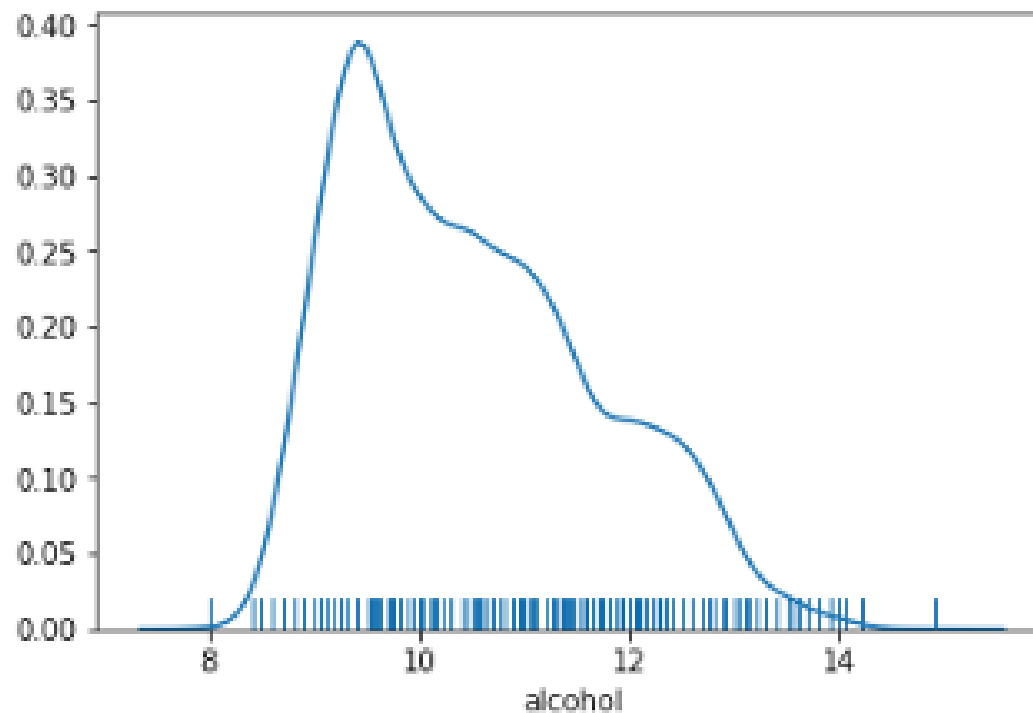
```
sns.distplot(df['alcohol'], kde=False, bins=10)
```



Alternative data distributions

- A rug plot is an alternative way to view the distribution of data
- A kde curve and rug plot can be combined


```
sns.distplot(df['alcohol'], hist=False, rug=True)
```



Further Customizations

- The `distplot` function uses several functions including `kdeplot` and `rugplot`
- It is possible to further customize a plot by passing arguments to the underlying function

```
sns.distplot(df['alcohol'], hist=False,  
             rug=True, kde_kws={'shade': True})
```



Let's practice!

INTERMEDIATE DATA VISUALIZATION WITH SEABORN

Regression Plots in Seaborn

INTERMEDIATE DATA VISUALIZATION WITH SEABORN

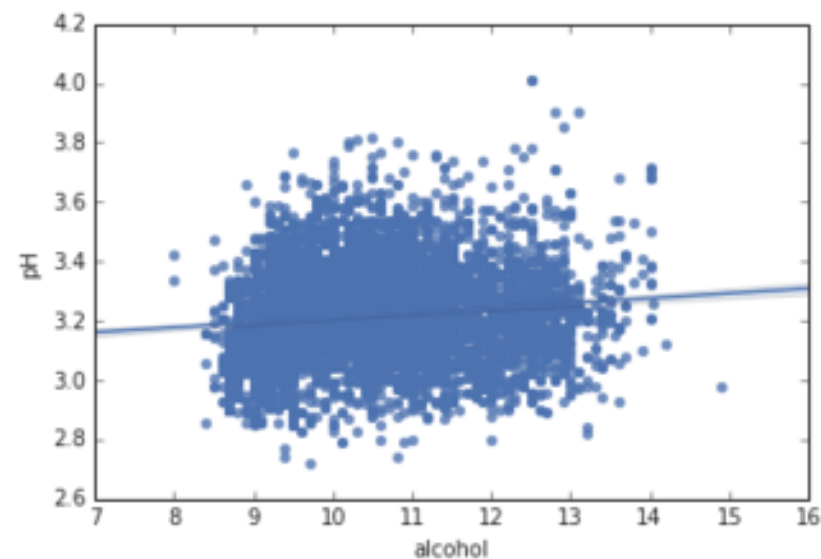


Chris Moffitt
Instructor

Introduction to regplot

- The `regplot` function generates a scatter plot with a regression line
- Usage is similar to the `distplot`
- The `data` and `x` and `y` variables must be defined

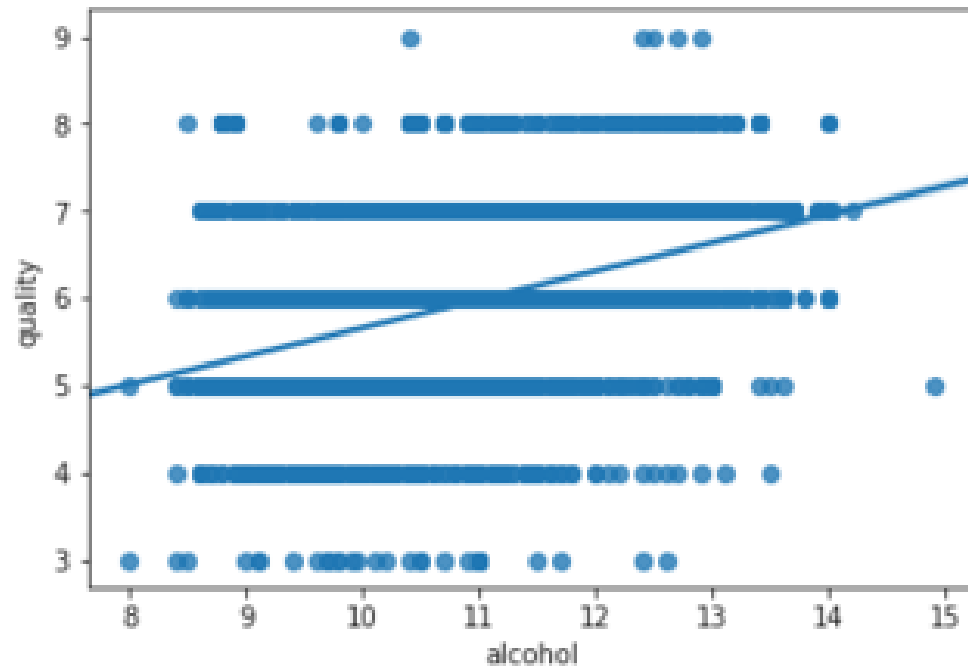
```
sns.regplot(x="alcohol", y="pH", data=df)
```



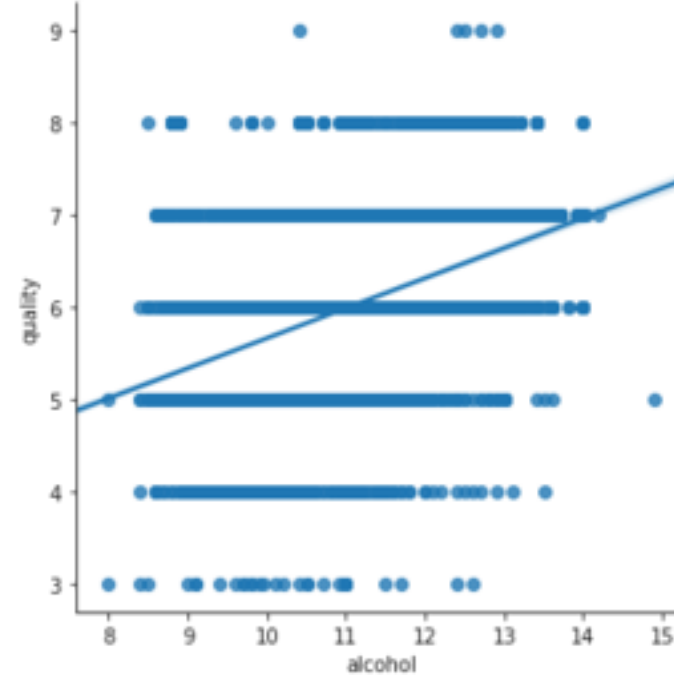
Implot() builds on top of the base regplot()

- `regplot` - low level
- `lplot` - high level

```
sns.regplot(x="alcohol",  
            y="quality",  
            data=df)
```



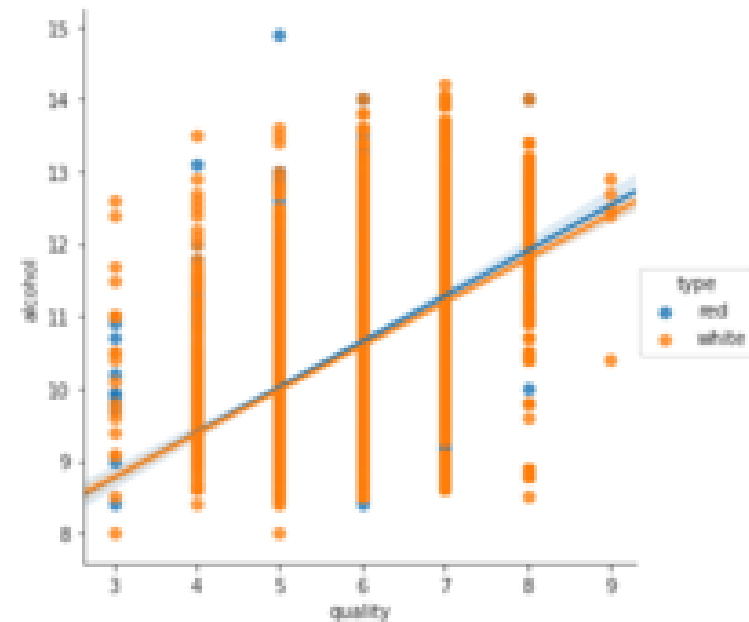
```
sns.lplot(x="alcohol",  
          y="quality",  
          data=df)
```



Implot faceting

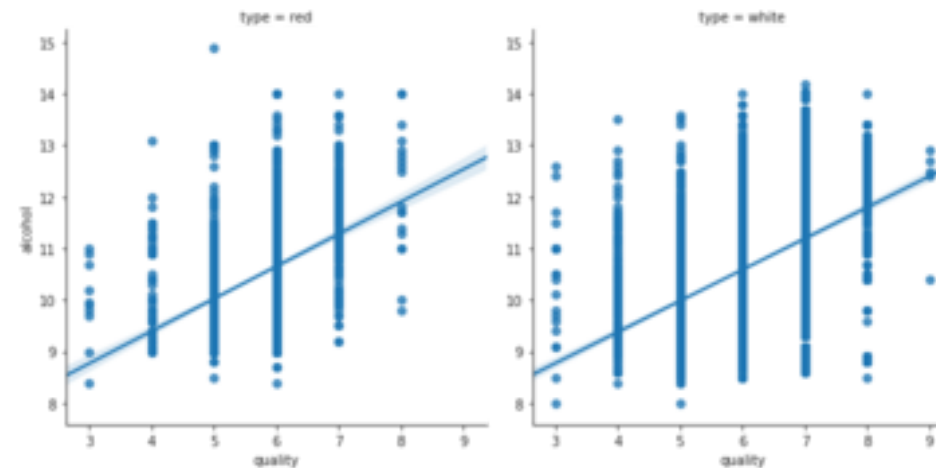
- Organize data by colors (`hue`)

```
sns.lmplot(x="quality",  
           y="alcohol",  
           data=df,  
           hue="type")
```



- Organize data by columns (`col`)

```
sns.lmplot(x="quality",  
           y="alcohol",  
           data=df,  
           col="type")
```



Let's practice!

INTERMEDIATE DATA VISUALIZATION WITH SEABORN

Using Seaborn Styles

INTERMEDIATE DATA VISUALIZATION WITH SEABORN




Chris Moffitt
Instructor

Setting Styles


- Seaborn has default configurations that can be applied with `sns.set()`
- These styles can override matplotlib and pandas plots as well

```
sns.set()  
df['Tuition'].plot.hist()
```

Pandas histogram




default



Seaborn style

Theme examples with sns.set_style()


```
for style in ['white', 'dark', 'whitegrid', 'darkgrid',  
             'ticks']:  
    sns.set_style(style)  
    sns.distplot(df['Tuition'])  
    plt.show()
```



Removing axes with `despine()`

- Sometimes plots are improved by removing elements
- Seaborn contains a shortcut for removing the spines of a plot

```
sns.set_style('white')  
sns.distplot(df['Tuition'])  
sns.despine(left=True)
```



Let's practice!

INTERMEDIATE DATA VISUALIZATION WITH SEABORN

Colors in Seaborn

INTERMEDIATE DATA VISUALIZATION WITH SEABORN




Chris Moffitt
Instructor

Defining a color for a plot

- Seaborn supports assigning colors to plots using `matplotlib` color codes


```
sns.set(color_codes=True)
sns.distplot(df['Tuition'], color='g')
```



Palettes

- Seaborn uses the `set_palette()` function to define a palette

```
for p in sns.palettes.SEABORN_PALETTES:  
    sns.set_palette(p)  
    sns.distplot(df['Tuition'])
```



Displaying Palettes

- `sns.palplot()` function displays a palette
- `sns.color_palette()` returns the current palette

```
for p in sns.palettes.SEABORN_PALETTES:  
    sns.set_palette(p)  
    sns.palplot(sns.color_palette())  
    plt.show()
```



Defining Custom Palettes

- Circular colors = when the data is not ordered

```
sns.palplot(sns.color_palette(  
    "Paired", 12))
```




- Sequential colors = when the data has a consistent range from high to low

```
sns.palplot(sns.color_palette(  
    "Blues", 12))
```



- Diverging colors = when both the low and high values are interesting

```
sns.palplot(sns.color_palette(  
    "BrBG", 12))
```



Let's practice!

INTERMEDIATE DATA VISUALIZATION WITH SEABORN

Customizing with matplotlib

INTERMEDIATE DATA VISUALIZATION WITH SEABORN




Chris Moffitt
Instructor

Matplotlib Axes

- Most customization available through `matplotlib` `Axes` objects
- `Axes` can be passed to seaborn functions


```
fig, ax = plt.subplots()
sns.distplot(df['Tuition'], ax=ax)
ax.set(xlabel="Tuition 2013-14")
```



Further Customizations

- The `axes` object supports many common customizations

```
fig, ax = plt.subplots()
sns.distplot(df['Tuition'], ax=ax)
ax.set(xlabel="Tuition 2013-14",
       ylabel="Distribution", xlim=(0, 50000),
       title="2013-14 Tuition and Fees Distribution")
```



Combining Plots



- It is possible to combine and configure multiple plots

```
fig, (ax0, ax1) = plt.subplots(
    nrows=1, ncols=2, sharey=True, figsize=(7,4))

sns.distplot(df['Tuition'], ax=ax0)
sns.distplot(df.query(
    'State == "MN"')['Tuition'], ax=ax1)

ax1.set(xlabel="Tuition (MN)", xlim=(0, 70000))
ax1.axvline(x=20000, label='My Budget', linestyle='--')
ax1.legend()
```

Combining Plots



Let's practice!

INTERMEDIATE DATA VISUALIZATION WITH SEABORN

Categorical Plot Types

INTERMEDIATE DATA VISUALIZATION WITH SEABORN





Chris Moffitt
Instructor




Categorical Data

- Data which takes on a limited and fixed number of values
- Normally combined with numeric data
- Examples include:
 - Geography (country, state, region)
 - Gender
 - Ethnicity
 - Blood type
 - Eye color

Plot types - show each observation




Plot types - abstract representations




Plot types - statistical estimates


barplot



pointplot




countplot




Plots of each observation - stripplot

```
sns.stripplot(data=df, y="DRG Definition",  
              x="Average Covered Charges",  
              jitter=True)
```




Plots of each observation - swarmplot

```
sns.swarmplot(data=df, y="DRG Definition",  
              x="Average Covered Charges")
```




Abstract representations - boxplot

```
sns.boxplot(data=df, y="DRG Definition",  
            x="Average Covered Charges")
```




Abstract representation - violinplot

```
sns.violinplot(data=df, y="DRG Definition",  
               x="Average Covered Charges")
```




Abstract representation - lvplot

```
sns.lvplot(data=df, y="DRG Definition",  
           x="Average Covered Charges")
```




Statistical estimates - barplot

```
sns.barplot(data=df, y="DRG Definition",  
            x="Average Covered Charges",  
            hue="Region")
```




Statistical estimates - pointplot

```
sns.pointplot(data=df, y="DRG Definition",  
              x="Average Covered Charges",  
              hue="Region")
```



Statistical estimates - countplot

```
sns.countplot(data=df, y="DRG_Code", hue="Region")
```



Let's practice!

INTERMEDIATE DATA VISUALIZATION WITH SEABORN

Regression Plots

INTERMEDIATE DATA VISUALIZATION WITH SEABORN




Chris Moffitt
Instructor

Bicycle Dataset

- Aggregated bicycle sharing data in Washington DC
- Data includes:
 - Rental amounts
 - Weather information
 - Calendar information
- Can we predict rental amounts?

Plotting with regplot()


```
sns.regplot(data=df, x='temp',  
            y='total_rentals', marker='+')
```



Evaluating regression with residplot()

- A residual plot is useful for evaluating the fit of a model
- Seaborn supports through `residplot` function


```
sns.residplot(data=df, x='temp', y='total_rentals')
```



Polynomial regression


- Seaborn supports polynomial regression using the `order` parameter

```
sns.regplot(data=df, x='temp',  
            y='total_rentals', order=2)
```




residplot with polynomial regression

```
sns.residplot(data=df, x='temp',  
              y='total_rentals', order=2)
```



Categorical values


```
sns.regplot(data=df, x='mnth', y='total_rentals',  
            x_jitter=.1, order=2)
```



Estimators

- In some cases, an `x_estimator` can be useful for highlighting trends


```
sns.regplot(data=df, x='mnth', y='total_rentals',  
            x_estimator=np.mean, order=2)
```



Binning the data

- `x_bins` can be used to divide the data into discrete bins
- The regression line is still fit against all the data

```
sns.regplot(data=df, x='temp', y='total_rentals',  
            x_bins=4)
```



Let's practice!

INTERMEDIATE DATA VISUALIZATION WITH SEABORN

Matrix Plots

INTERMEDIATE DATA VISUALIZATION WITH SEABORN



Chris Moffitt
Instructor

Getting data in the right format


- Seaborn's `heatmap()` function requires data to be in a grid format
- pandas `crosstab()` is frequently used to manipulate the data

```
pd.crosstab(df["mnth"], df["weekday"],  
values=df["total_rentals"],aggfunc='mean').round(0)
```

weekday	0	1	2	3	4	5	6
mnth							
1	1816.0	1927.0	2568.0	2139.0	2513.0	2446.0	1957.0
2	2248.0	2604.0	2824.0	2813.0	2878.0	2933.0	2266.0
3	3301.0	3546.0	3574.0	3670.0	3817.0	3926.0	3939.0
4	4417.0	4516.0	4556.0	4331.0	4764.0	4387.0	4446.0
5	5320.0	4512.0	5025.0	5119.0	5893.0	5751.0	5978.0
6	5940.0	5478.0	5681.0	5701.0	5622.0	5616.0	6344.0
7	5298.0	5792.0	5844.0	5814.0	5624.0	5406.0	5232.0
8	4703.0	5518.0	5930.0	6077.0	6038.0	5958.0	5224.0
9	6160.0	5637.0	5184.0	5668.0	5486.0	5747.0	6394.0
10	4735.0	4632.0	5065.0	5505.0	5537.0	5623.0	5445.0
11	4126.0	4658.0	4040.0	4136.0	3994.0	4524.0	4288.0
12	2740.0	3498.0	3713.0	3270.0	3711.0	3742.0	3195.0


Build a heatmap

```
sns.heatmap(pd.crosstab(df["mnth"], df["weekday"],  
                        values=df["total_rentals"], aggfunc='mean')  
)
```



Customize a heatmap


```
sns.heatmap(df_crosstab, annot=True, fmt="d",  
            cmap="YlGnBu", cbar=False, linewidths=.5)
```



Centering a heatmap

- Seaborn support centering the heatmap colors on a specific value


```
sns.heatmap(df_crosstab, annot=True, fmt="d",  
            cmap="YlGnBu", cbar=True,  
            center=df_crosstab.loc[9, 6])
```



Plotting a correlation matrix

- Pandas `corr` function calculates correlations between columns in a dataframe
- The output can be converted to a heatmap with seaborn

```
sns.heatmap(df.corr())
```



Let's practice!

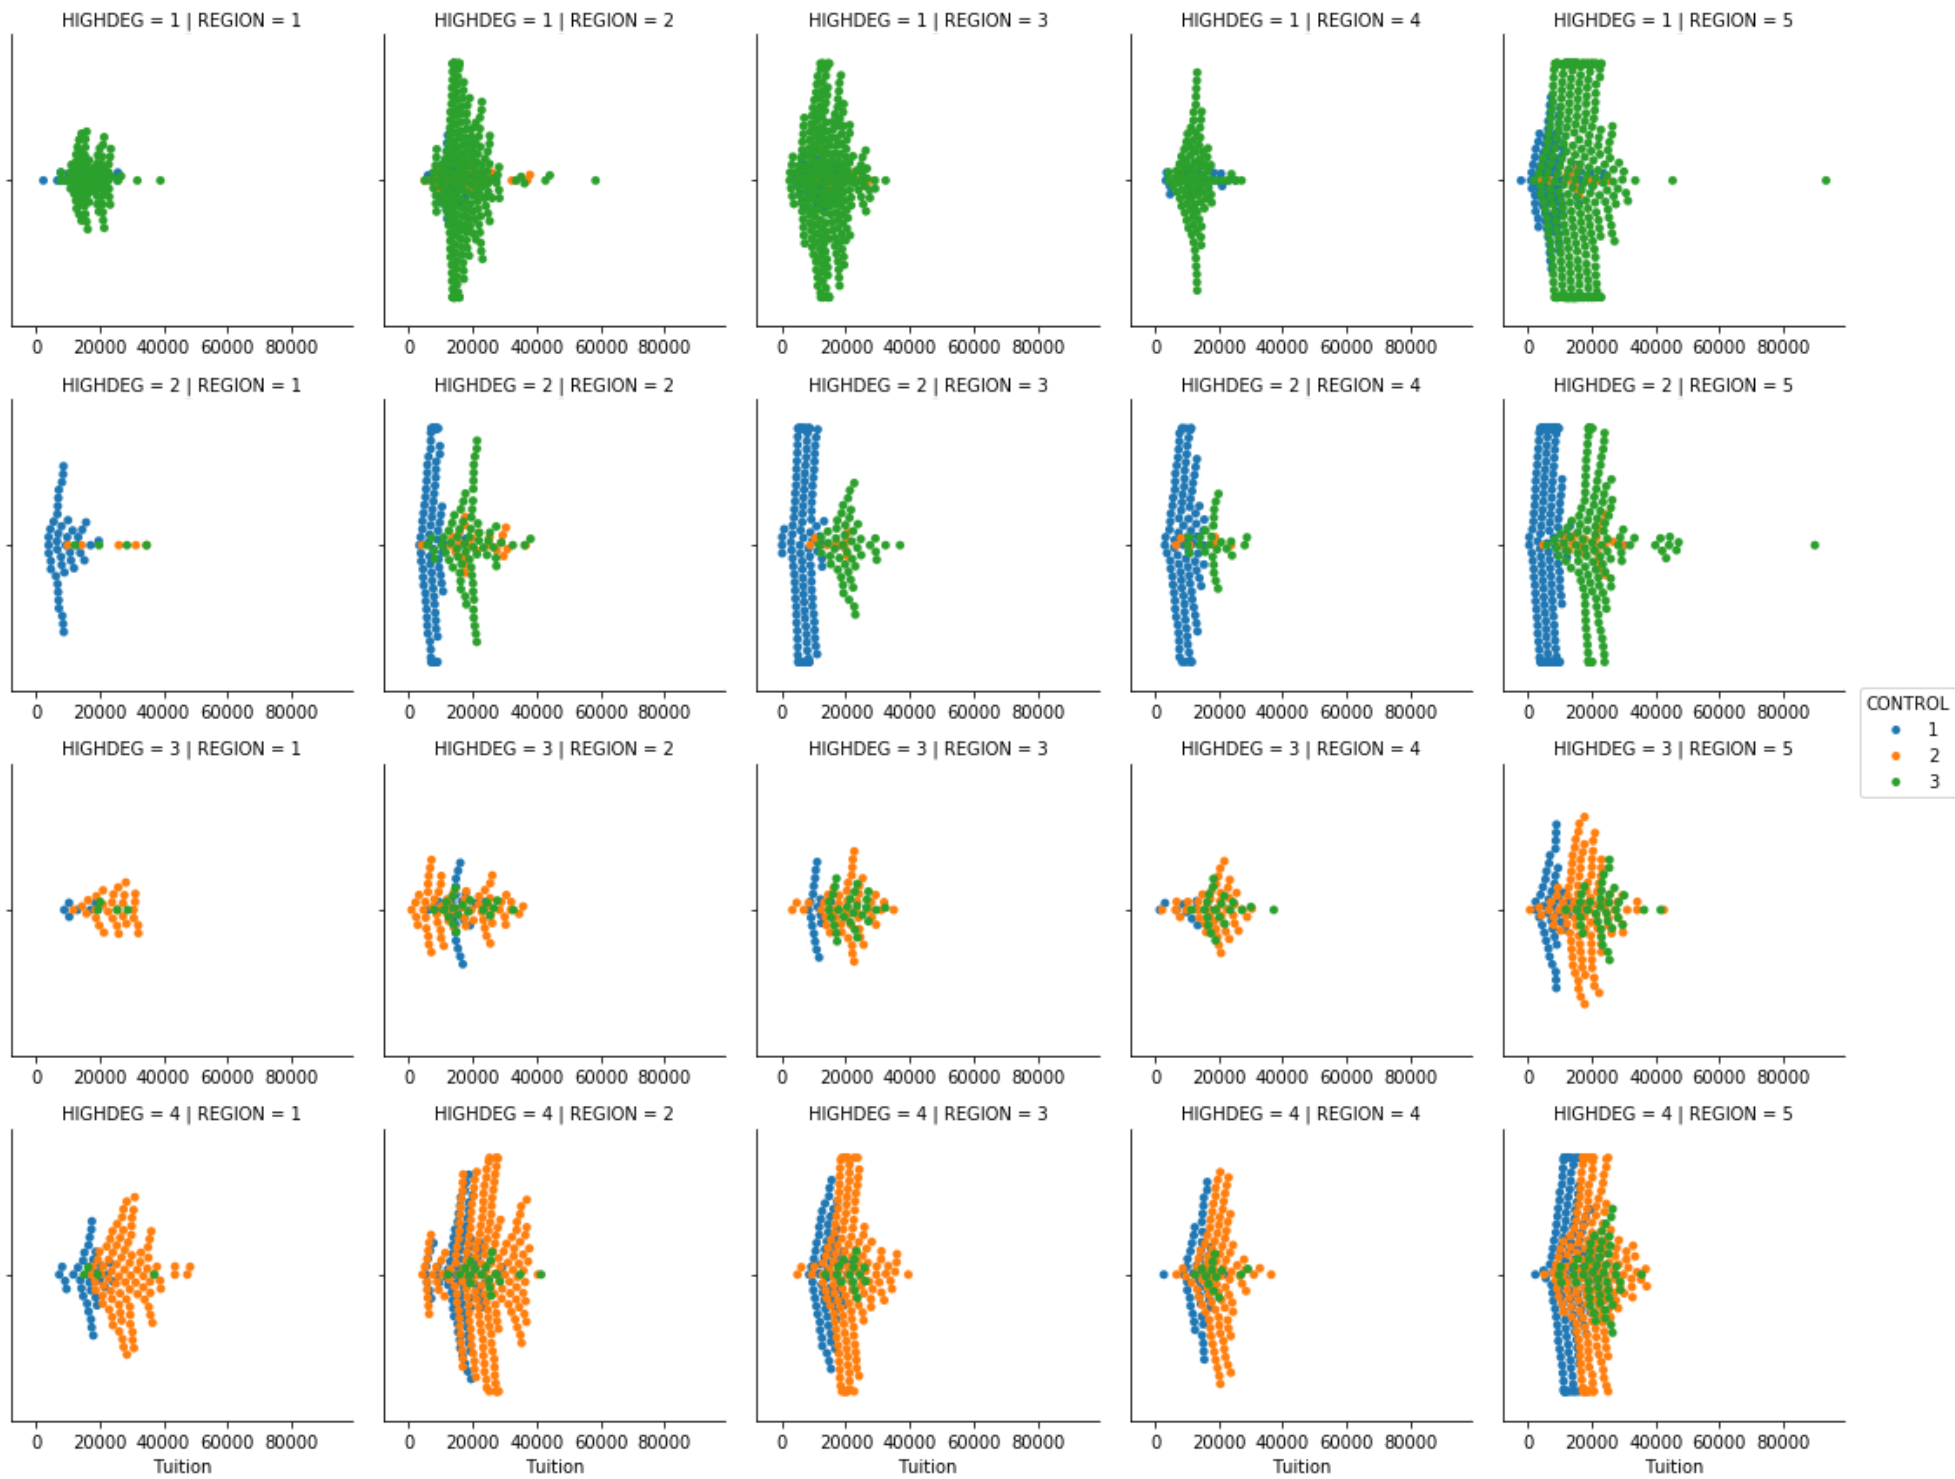
INTERMEDIATE DATA VISUALIZATION WITH SEABORN

Using FacetGrid, factorplot and Implot

INTERMEDIATE DATA VISUALIZATION WITH SEABORN



Chris Moffitt
Instructor



Tidy data

- Seaborn's grid plots require data in "tidy format"
- One observation per row of data


	INSTNM	OPEID	REGION	SAT_AVG_ALL	PCTPELL	PCTFLOAN	ADM_RATE_ALL	UG	AVGFACSAL	COMPL_RPY_5YR_RT	DEBT_MDN
0	Alabama A & M University	100200	5	850.0	0.7249	0.8159	0.653841	4380.0	7017.0	0.477631579	14600
1	University of Alabama at Birmingham	105200	5	1147.0	0.3505	0.5218	0.604275	10331.0	10221.0	0.673230442	14250
2	Amridge University	2503400	5	NaN	0.7455	0.8781	NaN	98.0	3217.0	0.636363636	11082
3	University of Alabama in Huntsville	105500	5	1221.0	0.3179	0.4589	0.811971	5220.0	9514.0	0.762222222	15000
4	Alabama State University	100500	5	844.0	0.7567	0.7692	0.463858	4348.0	7940.0	0.43006993	15274

FacetGrid

- The `FacetGrid` is foundational for many data aware grids
- It allows the user to control how data is distributed across columns, rows and hue
- Once a `FacetGrid` is created, the plot type must be mapped to the grid

FacetGrid Categorical Example


```
g = sns.FacetGrid(df, col="HIGHDEG")
g.map(sns.boxplot, 'Tuition',
      order=['1', '2', '3', '4'])
```



factorplot()

- The `factorplot` is a simpler way to use a `FacetGrid` for categorical data
- Combines the facetting and mapping process into 1 function


```
sns.factorplot(x="Tuition", data=df,  
               col="HIGHDEG", kind='box')
```



FacetGrid for regression

- `FacetGrid()` can also be used for scatter or regression plots


```
g = sns.FacetGrid(df, col="HIGHDEG")  
g.map(plt.scatter, 'Tuition', 'SAT_AVG_ALL')
```



Implot

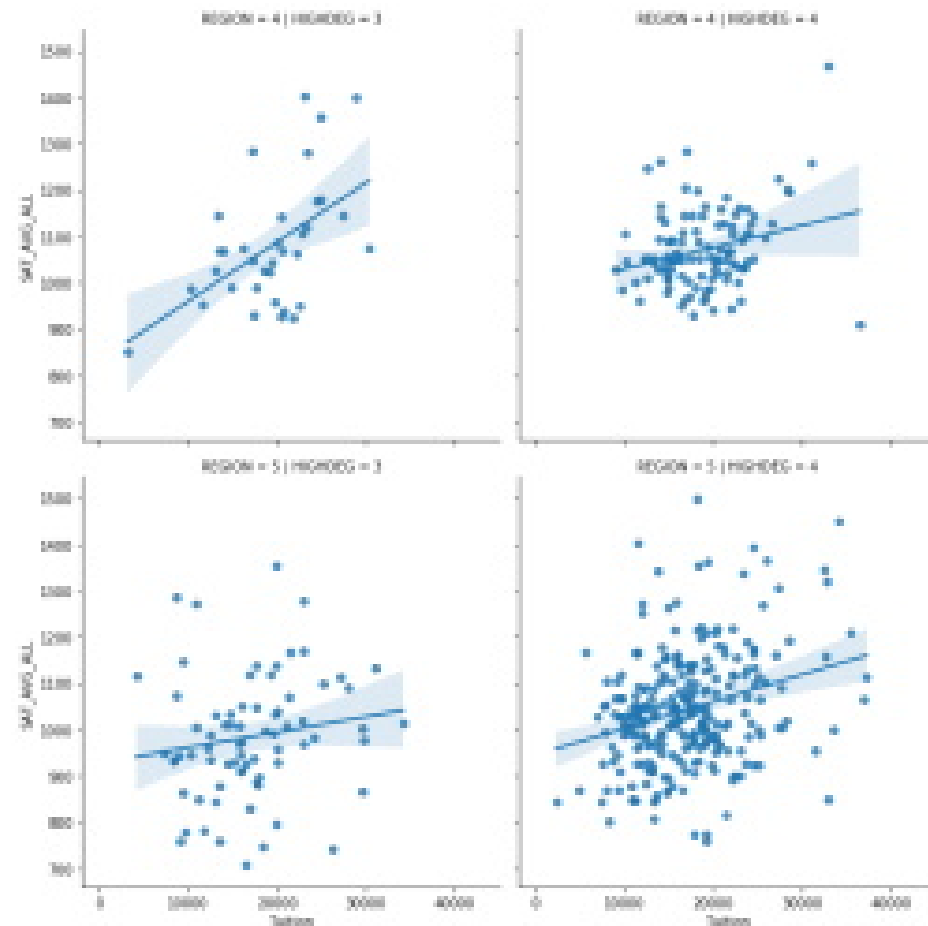
- `lmpplot` plots scatter and regression plots on a `FacetGrid`

```
sns.lmpplot(data=df, x="Tuition", y="SAT_AVG_ALL",  
            col="HIGHDEG", fit_reg=False)
```



Implot with regression

```
sns.lmplot(data=df, x="Tuition", y="SAT_AVG_ALL",  
           col="HIGHDEG", row='REGION')
```



Let's practice!

INTERMEDIATE DATA VISUALIZATION WITH SEABORN

Using PairGrid and pairplot


INTERMEDIATE DATA VISUALIZATION WITH SEABORN



Chris Moffitt
Instructor

Pairwise relationships


- `PairGrid` shows pairwise relationships between data elements



Creating a PairGrid


- The `PairGrid` follows similar API to `FacetGrid`

```
g = sns.PairGrid(df, vars=["Fair_Mrkt_Rent",  
                           "Median_Income"])  
  
g = g.map(plt.scatter)
```



Customizing the PairGrid diagonals


```
g = sns.PairGrid(df, vars=["Fair_Mrkt_Rent",  
                           "Median_Income"])  
  
g = g.map_diag(plt.hist)  
g = g.map_offdiag(plt.scatter)
```



Pairplot


- `pairplot` is a shortcut for the `PairGrid`

```
sns.pairplot(df, vars=["Fair_Mrkt_Rent",  
                      "Median_Income"], kind='reg',  
             diag_kind='hist')
```



Customizing a pairplot

```
sns.pairplot(df.query('BEDRMS < 3'),  
             vars=["Fair_Mrkt_Rent",  
                  "Median_Income", "UTILITY"],  
             hue='BEDRMS', palette='husl',  
             plot_kws={'alpha': 0.5})
```




Let's practice!

INTERMEDIATE DATA VISUALIZATION WITH SEABORN


Using JointGrid and jointplot

INTERMEDIATE DATA VISUALIZATION WITH SEABORN




Chris Moffitt
Instructor

JointGrid() Overview




Basic JointGrid

```
g = sns.JointGrid(data=df, x="Tuition",  
                  y="ADM_RATE_ALL")  
g.plot(sns.regplot, sns.distplot)
```




Advanced JointGrid

```
g = sns.JointGrid(data=df, x="Tuition",  
                  y="ADM_RATE_ALL")  
  
g = g.plot_joint(sns.kdeplot)  
g = g.plot_marginals(sns.kdeplot, shade=True)  
g = g.annotate(stats.pearsonr)
```



jointplot()


```
sns.jointplot(data=df, x="Tuition",  
              y="ADM_RATE_ALL", kind='hex')
```



Customizing a jointplot

```
g = (sns.jointplot(x="Tuition",
                  y="ADM_RATE_ALL", kind='scatter',
                  xlim=(0, 25000),
                  marginal_kws=dict(
                      bins=15, rug=True),
                  data=df.query('UG < 2500 &
                               Ownership == "Public"'))
     .plot_joint(sns.kdeplot))
```

Customizing a jointplot



Let's practice!


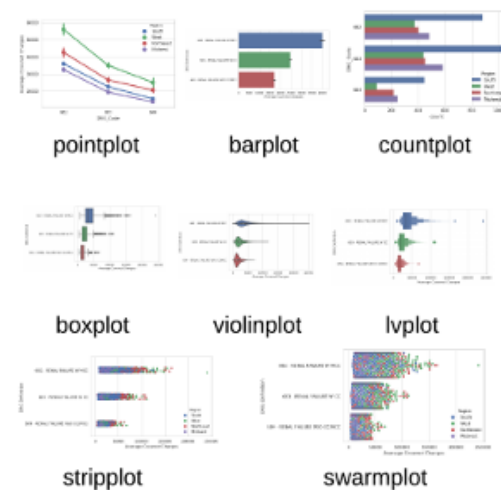
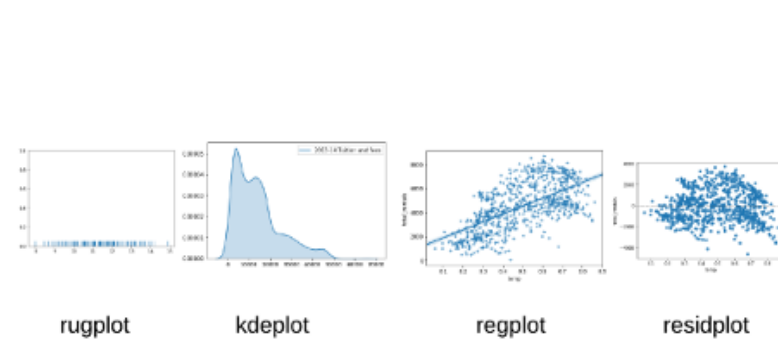



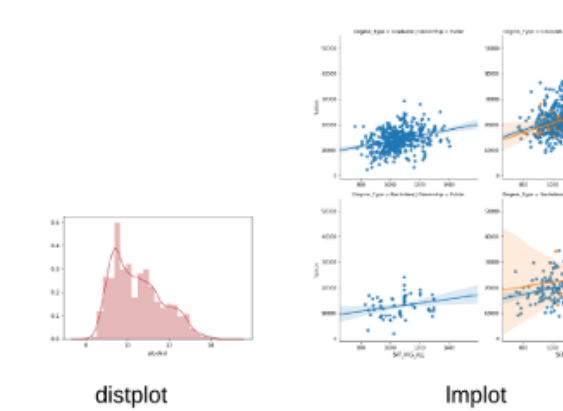


INTERMEDIATE DATA VISUALIZATION WITH SEABORN

Selecting Seaborn Plots

INTERMEDIATE DATA VISUALIZATION WITH SEABORN




Chris Moffitt
Instructor




Univariate Distribution Analysis


- `distplot()` is the best place to start for this analysis
- `rugplot()` and `kdeplot()` can be useful alternatives



distplot



rugplot




kdeplot

`plt.hist()`

matplotlib

Regression Analysis

- `lmplot()` performs regression analysis and supports facetting




lmplot

FacetGrid

Categorical Plots

- Explore data with the categorical plots and facet with




factorplot

FacetGrid




pairplot() and jointplot()

- Perform regression analysis with `lmplot`
- Analyze distributions with `distplot`



pairplot

PairGrid



jointplot

JointGrid

Thank You!

INTERMEDIATE DATA VISUALIZATION WITH SEABORN