

Welcome!

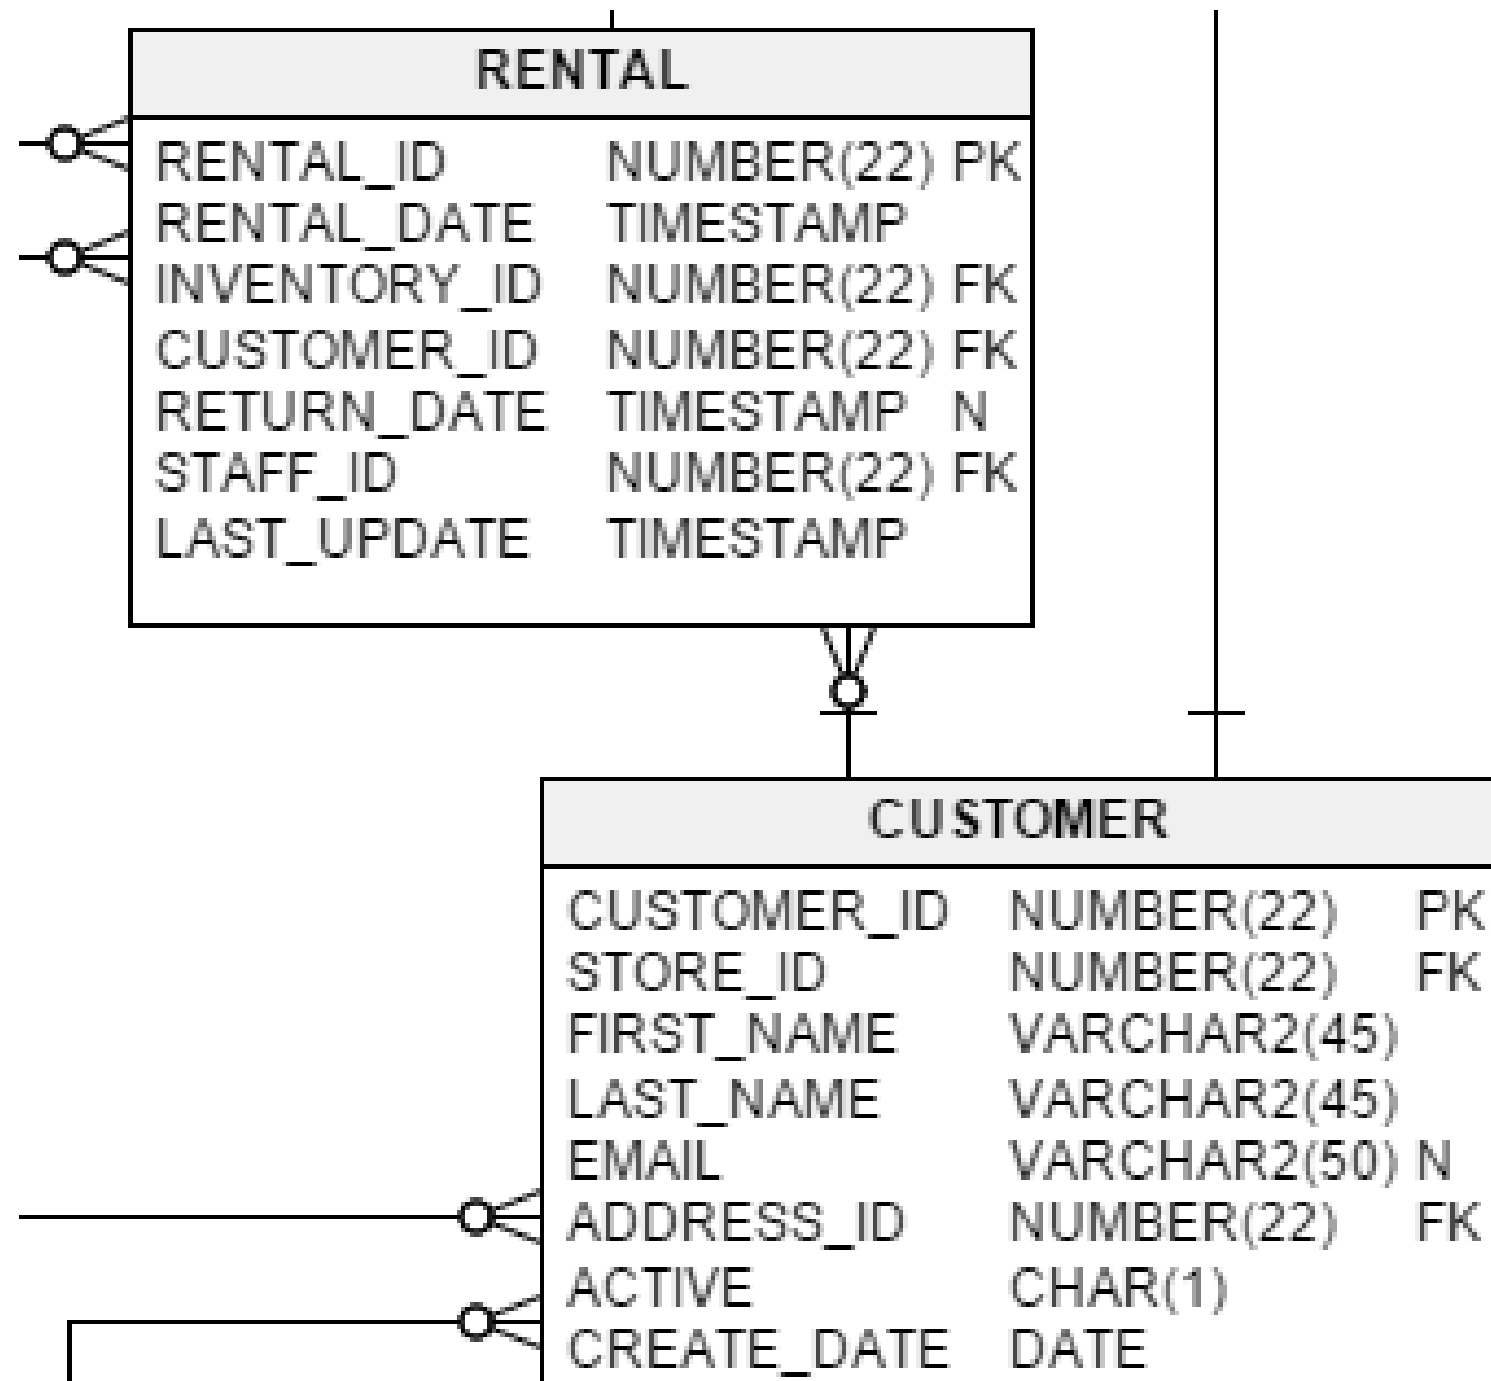
FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL



Brian Piccolo

Sr. Director, Digital Strategy

The Sakila Database



- Highly normalized
- Representative data types
- Custom functions

Topics

- Common data types in PostgreSQL
- Date and time functions and operators
- Parsing and manipulating text
- Full-text search and PostgreSQL Extensions

Common data types

- Text data types
 - `CHAR` , `VARCHAR` and `TEXT`
- Numeric data types
 - `INT` and `DECIMAL`
- Date / time data types
 - `DATE` , `TIME` , `TIMESTAMP` , `INTERVAL`
- Arrays

Text data types

```
SELECT title
FROM film
LIMIT 5
```

```
+-----+
| title          |
+-----+
| ACADEMY DINOSAUR |
| ACE GOLDFINGER  |
| ADAPTATION HOLES |
| AFFAIR PREJUDICE |
| AFRICAN EGG     |
+-----+
```

```
SELECT description
FROM film
LIMIT 2
```

```
+-----+
| description    |
+-----+
| A Epic Drama of a Feminist And a Mad |
| Scientist who must Battle a Teacher in |
| The Canadian Rockies.                 |
| A Astounding Epistle of a Database    |
| Administrator And a Explorer who     |
| must Find a Car in Ancient China      |
+-----+
```

Numeric data types

```
SELECT
```

```
    payment_id
```

```
FROM payment
```

```
LIMIT 5
```

```
+-----+
| payment_id |
+-----+
| 1          |
| 2          |
| 3          |
| 4          |
| 5          |
+-----+
```

```
SELECT
```

```
    amount
```

```
FROM payment
```

```
LIMIT 5
```

```
+-----+
| amount |
+-----+
| 2.99   |
| 0.99   |
| 5.99   |
| 0.99   |
| 9.99   |
+-----+
```

Determining data types from existing tables

```
SELECT
  title,
  description,
  special_features
FROM FILM
LIMIT 5
```

```
+-----+-----+-----+
| title      | description      | special_features      |
+-----+-----+-----+
| ACADEMY D... | A Epic...      | {Deleted Scenes,Behi...} |
| ACE GOLD...  | A Astound..    | {Trailers,Deleted Scenes} |
| AFFAIR PR... | A Fanciful,..  | {Commentaries,Behind the...} |
+-----+-----+-----+
```

Determining data types from existing tables

```
SELECT
    column_name,
    data_type
FROM INFORMATION_SCHEMA.COLUMNS
WHERE column_name in ('title', 'description', 'special_features')
AND table_name = 'film';
```

```
+-----+-----+
| column_name | data_type |
+-----+-----+
| title       | character varying |
| description | text          |
| special_features | ARRAY      |
+-----+-----+
```


Let's practice!

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

Date and time data types

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL



Brian Piccolo

Sr. Director, Digital Strategy

TIMESTAMP data types

- ISO 8601 format: yyyy-mm-dd

```
+-----+
| timestamp                |
|-----|
| 2019-03-26 01:05:17.93027+00 |
+-----+
```

```
SELECT payment_date
FROM payment;
```

```
+-----+
| payment_date            |
|-----|
| 2005-05-25 11:30:37    |
+-----+
```

DATE and TIME data types

```
+-----+-----+
| date       | time              |
+-----+-----+
| 2005-05-28 | 01:05:17.93027+00 |
+-----+-----+
```

```
SELECT create_date
FROM customer
```

```
+-----+
| create_date |
+-----+
| 2006-02-14   |
+-----+
```

INTERVAL data types

```
+-----+  
| interval |  
+-----+  
| 4 days   |  
+-----+
```

```
SELECT rental_date + INTERVAL '3 days' as expected_return  
FROM rental;
```

```
+-----+  
| expected_return |  
+-----+  
| 2005-05-27 22:53:30 |  
+-----+
```

Looking at date and time types

```
SELECT
    column_name,
    data_type
FROM INFORMATION_SCHEMA.COLUMNS
WHERE column_name in ('rental_date')
AND table_name = 'rental';
```

```
+-----+-----+
| column_name | data_type                |
|-----|-----|
| rental_date | timestamp without time zone |
+-----+-----+
```

Let's practice!

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

Working with ARRAYs

FUNCTIONS FOR MANIPULATING DATA IN POSTGRES SQL

SQL

Brian Piccolo

Sr. Director, Digital Strategy

Before we get started

CREATE TABLE example

```
CREATE TABLE my_first_table (  
  first_column text,  
  second_column integer  
);
```

INSERT example

```
INSERT INTO my_first_table  
  (first_column, second_column) VALUES ('text value', 12);
```

ARRAY a special type

Let's create a simple table with two array columns.

```
CREATE TABLE grades (  
  student_id int,  
  email text[],  
  test_scores int[]  
);
```

INSERT statements with ARRAYS

Example INSERT statement:

```
INSERT INTO grades  
VALUES (1,  
       '{"work","work1@datacamp.com"},"other","other1@datacamp.com"}',  
       '{92,85,96,88}' );
```

Accessing ARRAYS

```
SELECT
  email[1][1] AS type,
  email[1][2] AS address,
  test_scores[1],
FROM grades;
```

```
+-----+-----+-----+
| type   | address           | test_scores |
+-----+-----+-----+
| work   | work1@datacamp.com | 92          |
| work   | work2@datacamp.com | 76          |
+-----+-----+-----+
```

Note that PostgreSQL array indexes start with one and not zero.

Searching ARRAYS

SELECT

```
email[1][1] as type,  
email[1][2] as address,  
test_scores[1]
```

FROM grades

WHERE email[1][1] = 'work';

```
+-----+-----+-----+  
| type   | address                | test_scores |  
+-----+-----+-----+  
| work   | work1@datacamp.com     | 92          |  
| work   | work2@datacamp.com     | 76          |  
+-----+-----+-----+
```

ARRAY functions and operators

SELECT

```
email[2][1] as type,  
email[2][2] as address,  
test_scores[1]
```

FROM grades

WHERE 'other' = ANY (email);

```
+-----+-----+-----+  
| type   | address           | test_scores |  
+-----+-----+-----+  
| other  | other1@datacamp.com | 92          |  
| null   | null              | 76          |  
+-----+-----+-----+
```

ARRAY functions and operators

SELECT

```
email[2][1] as type,  
email[2][2] as address,  
test_scores[1]
```

FROM grades

WHERE email @> ARRAY['other'];

```
+-----+-----+-----+  
| type   | address           | test_scores |  
+-----+-----+-----+  
| other  | other1@datacamp.com | 92          |  
| null   | null              | 76          |  
+-----+-----+-----+
```

Let's practice!

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

Overview of basic arithmetic operators

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

SQL

Brian Piccolo

Sr. Director, Digital Strategy

Topics

- Overview of basic arithmetic operators
- The `CURRENT_DATE` , `CURRENT_TIMESTAMP` , `NOW()` functions
- The `AGE()` function
- The `EXTRACT()` , `DATE_PART()` , and `DATE_TRUNC()` functions

Adding and subtracting date / time data

```
SELECT date '2005-09-11' - date '2005-09-10';
```

```
+-----+  
| integer |  
+-----+  
| 1       |  
+-----+
```

Adding and subtracting date / time data

```
SELECT date '2005-09-11' + integer '3';
```

```
+-----+  
| date   |  
|-----|  
| 2005-09-14 |  
+-----+
```

Adding and subtracting date / time data

```
SELECT date '2005-09-11 00:00:00' - date '2005-09-09 12:00:00';
```

```
+-----+  
| interval |  
|-----|  
| 1 day 12:00:00 |  
+-----+
```

Calculating time periods with AGE

```
SELECT AGE(timestamp '2005-09-11 00:00:00', timestamp '2005-09-09 12:00:00');
```

```
+-----+  
| interval |  
|-----|  
| 1 day 12:00:00 |  
+-----+
```

DVDs, really??

```
SELECT
    AGE(rental_date)
FROM rental;
```

```
+-----+
| age                |
+-----+
| 13 years 11 mons 12 days 01:06:30 |
| 13 years 11 mons 12 days 01:05:27 |
| 13 years 11 mons 12 days 00:56:21 |
+-----+
```

Date / time arithmetic using INTERVALs

```
SELECT rental_date + INTERVAL '3 days' as expected_return  
FROM rental;
```

```
+-----+  
| expected_return |  
+-----+  
| 2005-05-27 22:53:30 |  
+-----+
```


Date / time arithmetic using INTERVALs

```
SELECT timestamp '2019-05-01' + 21 * INTERVAL '1 day';
```

```
+-----+  
| timestamp without timezone |  
+-----+  
| 2019-05-22 00:00:00        |  
+-----+
```

Let's practice!

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

Functions for retrieving current date/time

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

SQL

Brian Piccolo

Sr. Director, Digital Strategy

Retrieving the current timestamp

```
SELECT NOW();
```

```
+-----+  
| now() |  
+-----+  
| 2019-04-19 02:51:18.448641+00 |  
+-----+
```

Retrieving the current timestamp

```
SELECT NOW()::timestamp;
```

```
+-----+  
| now() |  
+-----+  
| 2019-04-19 02:51:18.448641 |  
+-----+
```

Retrieving the current timestamp

PostgreSQL specific casting

```
SELECT NOW()::timestamp;
```

CAST() function

```
SELECT CAST(NOW() as timestamp);
```

Retrieving the current timestamp

```
SELECT CURRENT_TIMESTAMP;
```

```
+-----+  
| current_timestamp |  
+-----+  
| 2019-04-19 02:51:18.448641+00 |  
+-----+
```

Retrieving the current timestamp

```
SELECT CURRENT_TIMESTAMP(2);
```

```
+-----+
| current_timestamp |
|-----|
| 2019-04-19 02:51:18.44+00 |
+-----+
```


Current date and time

```
SELECT CURRENT_DATE;
```

```
+-----+  
| current_date |  
|-----|  
| 2019-04-19   |  
+-----+
```

Current date and time

```
SELECT CURRENT_TIME;
```

```
+-----+  
| current_time |  
+-----+  
| 04:06:30.929845+00:00 |  
+-----+
```

Let's practice!

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

Extracting and transforming date / time data

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

SQL

Brian Piccolo

Sr. Director, Digital Strategy

Extracting and transforming date and time data

Exploring the `EXTRACT()`, `DATE_PART()` and `DATE_TRUNC()` functions

- Transactional timestamp precision not useful for analysis

```
2005-05-13 08:53:53
```

- Often need to extract parts of timestamps

```
2005 or 5 or 2 or Friday
```

- Or convert / truncate timestamp precision to standardize

```
2005-05-13 00:00:00
```

Extracting and transforming date / time data

- `EXTRACT(field FROM source)`

```
SELECT EXTRACT(quarter FROM timestamp '2005-01-24 05:12:00') AS quarter;
```

- `DATE_PART('field', source)`

```
SELECT DATE_PART('quarter', timestamp '2005-01-24 05:12:00') AS quarter;
```

```
+-----+  
| quarter |  
+-----+  
| 1       |  
+-----+
```

Extracting sub-fields from timestamp data

Transactional data from DVD Rentals *payment* table

```
SELECT * FROM payment;
```

```
+-----+-----+-----+-----+-----+-----+
| payment_id | customer_id | staff_id | rental_id | amount | payment_date |
|-----|-----|-----|-----|-----|-----|
| 1 | 1 | 1 | 76 | 2.99 | 2005-05-25 11:30:37 |
| 2 | 1 | 1 | 573 | 0.99 | 2005-05-28 10:35:23 |
| 3 | 1 | 1 | 1185 | 5.99 | 2005-06-15 0:54:12 |
+-----+-----+-----+-----+-----+-----+

```

Extracting sub-fields from timestamp data

Data from *payment* table by year and quarter Results

```
SELECT
  EXTRACT(quarter FROM payment_date) AS quarter,
  EXTRACT(year FROM payment_date) AS year,
  SUM(amount) AS total_payments
FROM
  payment
GROUP BY 1, 2;
```

```
+-----+
| quarter | year | total_payments |
+-----+-----+-----+
| 2       | 2005 | 14456.31      |
| 3       | 2005 | 52446.02      |
| 1       | 2006 | 514.18        |
+-----+-----+-----+
```


Truncating timestamps using DATE_TRUNC()

The `DATE_TRUNC()` function will truncate timestamp or interval data types.

- Truncate timestamp '2005-05-21 15:30:30' by year

```
SELECT DATE_TRUNC('year', TIMESTAMP '2005-05-21 15:30:30');
```

```
Result: 2005-01-01 00:00:00
```

- Truncate timestamp '2005-05-21 15:30:30' by month

```
SELECT DATE_TRUNC('month', TIMESTAMP '2005-05-21 15:30:30');
```

```
Result: 2005-05-01 00:00:00
```

Let's practice!

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

Reformatting string and character data

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

SQL

Brian Piccolo

Sr. Director, Digital Strategy

Topics

- Reformatting string and character data.
- Parsing string and character data.
- Determine string length and character position.
- Truncating and padding string data.

The string concatenation operator

```
SELECT
  first_name,
  last_name,
  first_name || ' ' || last_name AS full_name
FROM customer
```

```
+-----+-----+-----+
| first_name | last_name | full_name      |
|-----|-----|-----|
| MARY       | SMITH     | MARY SMITH     |
| LINDA      | WILLIAMS  | LINDA WILLIAMS  |
+-----+-----+-----+
```

String concatenation with functions

```
SELECT
  CONCAT(first_name, ' ', last_name) AS full_name
FROM customer;
```

```
+-----+
| first_name | last_name | full_name      |
+-----+
| MARY       | SMITH     | MARY SMITH     |
| LINDA      | WILLIAMS  | LINDA WILLIAMS |
+-----+
```

String concatenation with a non-string input

```
SELECT
  customer_id || ': '
  || first_name || ' '
  || last_name AS full_name
FROM customer;
```

```
+-----+
| full_name          |
+-----+
| 1: MARY SMITH      |
| 2: LINDA WILLIAMS  |
+-----+
```

Changing the case of string

```
SELECT
  UPPER(email)
FROM customer;
```

```
+-----+
| UPPER(email) |
+-----+
| MARY.SMITH@SAKILACUSTOMER.ORG |
| PATRICIA.JOHNSON@SAKILACUSTOMER.ORG |
| LINDA.WILLIAMS@SAKILACUSTOMER.ORG |
+-----+
```


Changing the case of string

```
SELECT  
  LOWER(title)  
FROM film;
```

```
+-----+  
| LOWER(title) |  
+-----+  
| academy dinosaur |  
| ace goldfinger |  
| adaptation holes |  
+-----+
```

Changing the case of string

```
SELECT  
  INITCAP(title)  
FROM film;
```

```
+-----+  
| INITCAP(title) |  
+-----+  
| Academy Dinosaur |  
| Ace Goldfinger   |  
| Adaptation Holes |  
+-----+
```

Replacing characters in a string

```
SELECT description FROM film;
```

```
+-----+
| description |
+-----+
| A Epic Drama of a Feminist And a Mad Scientist... |
| A Astounding Epistle of a Database Administrator... |
| A Astounding Reflection of a Lumberjack And a Car... |
| A Fanciful Documentary of a Frisbee And a Lumberjack... |
| A Fast-Paced Documentary of a Pastry Chef And a... |
+-----+
```

Replacing characters in a string

```
SELECT
    REPLACE(description, 'A Astounding',
             'An Astounding') as description
FROM film;
```

```
+-----+
| description                                     |
|-----|
| A Epic Drama of a Feminist And a Mad Scientist... |
| An Astounding Epistle of a Database Administrator... |
| An Astounding Reflection of a Lumberjack And a Car... |
+-----+
```

Manipulating string data with REVERSE

```
SELECT
  title,
  REVERSE(title)
FROM
  film AS f;
```

```
+-----+
| title          | reverse(title) |
+-----+
| ACADEMY DINOSAUR | RUASONID YMEDACA |
| ACE GOLDFINGER   | REGNIFDLOG ECA   |
+-----+
```

Let's practice!

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

Parsing string and character data

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

SQL

Brian Piccolo

Sr. Director, Digital Strategy

Determining the length of a string

```
SELECT
    title,
    CHAR_LENGTH(title)
FROM film;
```

```
+-----+-----+
| title          | CHAR_LENGTH(title) |
+-----+-----+
| ACADEMY DINOSAUR | 16                  |
| ACE GOLDFINGER   | 14                  |
| ADAPTATION HOLES | 16                  |
+-----+-----+
```


Determining the length of a string

```
SELECT
    title,
    LENGTH(title)
FROM film;
```

```
+-----+-----+
| title          | LENGTH(title) |
+-----+-----+
| ACADEMY DINOSAUR | 16            |
| ACE GOLDFINGER  | 14            |
| ADAPTATION HOLES | 16            |
+-----+-----+
```

Finding the position of a character in a string

```
SELECT
    email,
    POSITION('@' IN email)
FROM customer;
```

```
+-----+-----+
| email                                | POSITION('@' IN email) |
+-----+-----+
| MARY.SMITH@sakilacustomer.org       | 11                    |
| PATRICIA.JOHNSON@sakilacustomer.org | 17                    |
| LINDA.WILLIAMS@sakilacustomer.org   | 15                    |
+-----+-----+
```

Finding the position of a character in a string

```
SELECT
    email,
    STRPOS(email, '@')
FROM customer;
```

```
+-----+-----+
| email                                | STRPOS(email, '@') |
+-----+-----+
| MARY.SMITH@sakilacustomer.org        | 11                  |
| PATRICIA.JOHNSON@sakilacustomer.org  | 17                  |
| LINDA.WILLIAMS@sakilacustomer.org    | 15                  |
+-----+-----+
```

Parsing string data

```
SELECT
    LEFT(description, 50)
FROM film;
```

```
+-----+
| description |
+-----+
| A Epic Drama of a Feminist And a Mad Scientist who |
| A Astounding Epistle of a Database Administrator A |
| A Astounding Reflection of a Lumberjack And a Car  |
+-----+
```

Parsing string data

```
SELECT
    RIGHT(description, 50)
FROM film;
```

```
+-----+
| description |
+-----+
|  who must Battle a Teacher in The Canadian Rockies |
| nd a Explorer who must Find a Car in Ancient China |
| Car who must Sink a Lumberjack in A Baloon Factory |
+-----+
```

Extracting substrings of character data

```
SELECT
    SUBSTRING(description, 10, 50)
FROM
    film AS f;
```

```
+-----+
| description |
+-----+
| ama of a Feminist And a Mad Scientist who must Bat |
| ing Epistle of a Database Administrator And a Expl |
| ing Reflection of a Lumberjack And a Car who must |
+-----+
```

Extracting substrings of character data

```
SELECT
    SUBSTRING(email FROM 0 FOR POSITION('@' IN email))
FROM
    customer;
```

```
+-----+
| SUBSTRING(email FROM 0 FOR POSITION('@' IN email)) |
|-----|
| MARY.SMITH                                         |
| PATRICIA.JOHNSON                                  |
| LINDA.WILLIAMS                                     |
+-----+
```

Extracting substrings of character data

```
SELECT
    SUBSTRING(email FROM POSITION('@' IN email)+1 FOR CHAR_LENGTH(email))
FROM
    customer;
```

```
+-----+
| SUBSTRING(email FROM POSITION('@' IN email)+1 FOR CHAR_LENGTH(email)) |
|-----|
| sakilacustomer.org |
| sakilacustomer.org |
| sakilacustomer.org |
+-----+
```


Extracting substrings of character data

```
SELECT
    SUBSTR(description, 10, 50)
FROM
    film AS f;
```

```
+-----+
| description |
+-----+
| ama of a Feminist And a Mad Scientist who must Bat |
| ing Epistle of a Database Administrator And a Expl |
| ing Reflection of a Lumberjack And a Car who must |
+-----+
```

Let's practice!

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

Truncating and padding string data

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

SQL

Brian Piccolo

Sr. Director, Digital Strategy

Removing whitespace from strings

```
TRIM([leading | trailing | both] [characters] from string)
```

- **First parameter:** [leading | trailing | both]
- **Second parameter:** [characters]
- **Third parameter:** from string

Removing whitespace from strings

```
SELECT TRIM(' padded ');
```

```
+-----+  
| TRIM   |  
|-----|  
| padded |  
+-----+
```

Removing whitespace from strings

```
SELECT LTRIM(' padded');
```

```
+-----+  
| LTRIM |  
|-----|  
| padded |  
+-----+
```

Removing whitespace from strings

```
SELECT RTRIM(' padded');
```

```
+-----+  
| RTRIM |  
+-----+  
| padded |  
+-----+
```

Padding strings with character data

```
SELECT LPAD('padded', 10, '#');
```

```
+-----+  
| LPAD   |  
|-----|  
| ####padded |  
+-----+
```


Padding strings with whitespace

```
SELECT LPAD('padded', 10);
```

```
+-----+
| LPAD   |
|-----|
| padded |
+-----+
```

```
SELECT LPAD('padded', 5);
```

```
+-----+
| LPAD   |
|-----|
| padde  |
+-----+
```

Padding strings with whitespace

```
SELECT RPAD('padded', 10, '#');
```

```
+-----+  
| RPAD   |  
|-----|  
| padded#### |  
+-----+
```

Let's practice!

FUNCTIONS FOR MANIPULATING DATA IN POSTGRESQL

Pivoting

POSTGRESQL SUMMARY STATS AND WINDOW FUNCTIONS



Michel Semaan
Data Scientist

Transforming tables

Before

Country	Year	Awards
CHN	2008	74
CHN	2012	56
RUS	2008	43
RUS	2012	47
USA	2008	125
USA	2012	147

- Gold medals awarded to China, Russia, and the USA

After

Country	2008	2012
CHN	74	56
RUS	43	47
USA	125	147

- Pivoted by Year
- Easier to scan, especially if pivoted by a chronologically ordered column

Enter CROSSTAB

```
CREATE EXTENSION IF NOT EXISTS tablefunc;
```

```
SELECT * FROM CROSSTAB($$  
  source_sql TEXT  
$$) AS ct (column_1 DATA_TYPE_1,  
           column_2 DATA_TYPE_2,  
           ...,  
           column_n DATA_TYPE_N);
```

Queries

Before

```
SELECT
  Country, Year, COUNT(*) AS Awards
FROM Summer_Medals
WHERE
  Country IN ('CHN', 'RUS', 'USA')
  AND Year IN (2008, 2012)
  AND Medal = 'Gold'
GROUP BY Country, Year
ORDER BY Country ASC, Year ASC;
```

After

```
CREATE EXTENSION IF NOT EXISTS tablefunc;

SELECT * FROM CROSSTAB($$
  SELECT
    Country, Year, COUNT(*) :: INTEGER AS Awards
  FROM Summer_Medals
  WHERE
    Country IN ('CHN', 'RUS', 'USA')
    AND Year IN (2008, 2012)
    AND Medal = 'Gold'
  GROUP BY Country, Year
  ORDER BY Country ASC, Year ASC;
$$) AS ct (Country VARCHAR, "2008" INTEGER, "2012" INTEGER)
ORDER BY Country ASC;
```

Source query

```
WITH Country_Awards AS (  
  SELECT  
    Country, Year, COUNT(*) AS Awards  
  FROM Summer_Medals  
  WHERE  
    Country IN ('CHN', 'RUS', 'USA')  
    AND Year IN (2004, 2008, 2012)  
    AND Medal = 'Gold' AND Sport = 'Gymnastics'  
  GROUP BY Country, Year  
  ORDER BY Country ASC, Year ASC)  
  
SELECT  
  Country, Year,  
  RANK() OVER  
    (PARTITION BY Year ORDER BY Awards DESC) :: INTEGER  
  AS rank  
FROM Country_Awards  
ORDER BY Country ASC, Year ASC;
```


Source result

Country	Year	Rank
CHN	2004	3
CHN	2008	1
CHN	2012	1
RUS	2004	1
RUS	2008	2
RUS	2012	2
USA	2004	2
USA	2008	3
USA	2012	3

Pivot query

```
CREATE EXTENSION IF NOT EXISTS tablefunc;
```

```
SELECT * FROM CROSSTAB($$  
  ...  
$$) AS ct (Country VARCHAR,  
           "2004" INTEGER,  
           "2008" INTEGER,  
           "2012" INTEGER)
```

```
ORDER BY Country ASC;
```

Pivot result

Country	2004	2008	2012
CHN	3	1	1
RUS	1	2	2
USA	2	3	3

Let's practice!

POSTGRESQL SUMMARY STATS AND WINDOW FUNCTIONS

ROLLUP and CUBE

POSTGRESQL SUMMARY STATS AND WINDOW FUNCTIONS



Michel Semaan
Data Scientist

Group-level totals

Chinese and Russian medals in the 2008 Summer Olympics per medal class

Country	Medal	Awards
CHN	Bronze	57
CHN	Gold	74
CHN	Silver	53
CHN	Total	184
RUS	Bronze	56
RUS	Gold	43
RUS	Silver	44
RUS	Total	143

The old way

```
SELECT
  Country, Medal, COUNT(*) AS Awards
FROM Summer_Medals
WHERE
  Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY Country, Medal
ORDER BY Country ASC, Medal ASC
UNION ALL
```

```
SELECT
  Country, 'Total', COUNT(*) AS Awards
FROM Summer_Medals
WHERE
  Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY Country, 2
ORDER BY Country ASC;
```

Enter ROLLUP

```
SELECT
  Country, Medal, COUNT(*) AS Awards
FROM Summer_Medals
WHERE
  Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY Country, ROLLUP(Medal)
ORDER BY Country ASC, Medal ASC;
```

- `ROLLUP` is a `GROUP BY` subclause that includes extra rows for group-level aggregations
- `GROUP BY Country, ROLLUP(Medal)` will count all `Country` - and `Medal` -level totals, then count only `Country` -level totals and fill in `Medal` with `null` s for these rows

ROLLUP - Query

```
SELECT
  Country, Medal, COUNT(*) AS Awards
FROM summer_medals
WHERE
  Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY ROLLUP(Country, Medal)
ORDER BY Country ASC, Medal ASC;
```

- **ROLLUP** is hierarchical, de-aggregating from the leftmost provided column to the right-most
 - **ROLLUP(Country, Medal)** includes **Country** -level totals
 - **ROLLUP(Medal, Country)** includes **Medal** -level totals
 - Both include grand totals

ROLLUP - Result

Country	Medal	Awards
CHN	Bronze	57
CHN	Gold	74
CHN	Silver	53
CHN	null	184
RUS	Bronze	56
RUS	Gold	43
RUS	Silver	44
RUS	null	143
null	null	327

- Group-level totals contain `nulls` ; the row with all `null` s is the grand total
- Notice that it didn't include `Medal` -level totals, since it's `ROLLUP(Country, Medal)` and not `ROLLUP(Medal, Country)`

Enter CUBE

```
SELECT
  Country, Medal, COUNT(*) AS Awards
FROM summer_medals
WHERE
  Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY CUBE(Country, Medal)
ORDER BY Country ASC, Medal ASC;
```

- CUBE is a non-hierarchical ROLLUP
- It generates all possible group-level aggregations
 - CUBE(Country, Medal) counts Country -level, Medal -level, and grand totals

CUBE - Result

Country	Medal	Awards
CHN	Bronze	57
CHN	Gold	74
CHN	Silver	53
CHN	null	184
RUS	Bronze	56
RUS	Gold	43
RUS	Silver	44
RUS	null	143
null	Bronze	113
null	Gold	117
null	Silver	97
null	null	327

- Notice that `Medal` -level totals are included

ROLLUP vs CUBE

Source

Year	Quarter	Sales
2008	Q1	12
2008	Q2	15
2009	Q1	21
2009	Q2	27

- Use **ROLLUP** when you have hierarchical data (e.g., date parts) and don't want all possible group-level aggregations
- Use **CUBE** when you want all possible group-level aggregations

ROLLUP(Year, Quarter)

Year	Quarter	Sales
2008	null	27
2009	null	48
null	null	75

CUBE(Year, Quarter)

Above rows + the following

Year	Quarter	Sales
null	Q1	33
null	Q2	42

Let's practice!

POSTGRESQL SUMMARY STATS AND WINDOW FUNCTIONS

A survey of useful functions

POSTGRESQL SUMMARY STATS AND WINDOW FUNCTIONS



Michel Semaan
Data Scientist

Nulls ahoy

Query

```
SELECT
  Country, Medal, COUNT(*) AS Awards
FROM summer_medals
WHERE
  Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY ROLLUP(Country, Medal)
ORDER BY Country ASC, Medal ASC;
```

- null s signify group totals

Result

Country	Medal	Awards
CHN	Bronze	57
CHN	Gold	74
CHN	Silver	53
CHN	null	184
RUS	Bronze	56
RUS	Gold	43
RUS	Silver	44
RUS	null	143
null	null	327

Enter COALESCE

- `COALESCE()` takes a list of values and returns the first non-`null` value, going from left to right
- `COALESCE(null, null, 1, null, 2) ? 1`
- Useful when using SQL operations that return `null` s
 - `ROLLUP` and `CUBE`
 - Pivoting
 - `LAG` and `LEAD`

Annihilating nulls

Query

```
SELECT
  COALESCE(Country, 'Both countries') AS Country,
  COALESCE(Medal, 'All medals') AS Medal,
  COUNT(*) AS Awards
FROM summer_medals
WHERE
  Year = 2008 AND Country IN ('CHN', 'RUS')
GROUP BY ROLLUP(Country, Medal)
ORDER BY Country ASC, Medal ASC;
```

Result

Country	Medal	Awards
Both countries	All medals	327
CHN	All medals	184
CHN	Bronze	57
CHN	Gold	74
CHN	Silver	53
RUS	All medals	143
RUS	Bronze	56
RUS	Gold	43
RUS	Silver	44

Compressing data

Before

Country	Rank
CHN	1
RUS	2
USA	3

- Rank is redundant because the ranking is implied

After

CHN, RUS, USA

- Succinct and provides all information needed because the ranking is implied

Enter STRING_AGG

- `STRING_AGG(column, separator)` takes all the values of a column and concatenates them, with `separator` in between each value

`STRING_AGG(Letter, ', ')` transforms this...

```
| Letter |  
|-----|  
| A      |  
| B      |  
| C      |
```

...into this

```
A, B, C
```

Query and result

Before

```
WITH Country_Medals AS (  
  SELECT  
    Country, COUNT(*) AS Medals  
  FROM Summer_Medals  
  WHERE Year = 2012  
    AND Country IN ('CHN', 'RUS', 'USA')  
    AND Medal = 'Gold'  
    AND Sport = 'Gymnastics'  
  GROUP BY Country),  
  
  SELECT  
    Country,  
    RANK() OVER (ORDER BY Medals DESC) AS Rank  
  FROM Country_Medals  
  ORDER BY Rank ASC;
```

After

```
WITH Country_Medals AS (...),  
  
  Country_Ranks AS (...)  
  
  SELECT STRING_AGG(Country, ', ')  
  FROM Country_Medals;
```

Result

```
CHN, RUS, USA
```

Let's practice!

POSTGRESQL SUMMARY STATS AND WINDOW FUNCTIONS